Möbii³us Skill-Integration Audit — Grammatically Diagrammed Repository

`git4emet/doinbetter` :: `audits/skills/2026-05-skill-quad-parse.md`

> **TL;DR**
>
> - Three of four skills are canonical Anthropic open-source artifacts (`algorithmic-art`, `mcp-builder`, `skill-creator`, all Apache-2.0 from `anthropics/skills`); the fourth, `mobius-vocabulary`, is a Buddy-authored DoinBetter Collective skill with no public footprint and must be treated as the **integration substrate** that translates the other three into the 42-domain lattice.
> - The four skills compose cleanly into a single verb-pipeline — `mobius-vocabulary` GOVERNs and AUTHENTICATEs vocabulary; `skill-creator` MINTs and COORDINATEs new skills; `mcp-builder` SPEAKs/LISTENs across services; `algorithmic-art` RENDERs/VISUALIZEs — with no irreconcilable conflicts but three live decision points (template-branding override, evaluation gate adoption, "pushy description" voice).
> - Recommended commit path is staged: ship the four `shares.md/specs.md/questions.md/commentary.md` quartets immediately as audit artifacts, then resolve the three open questions before wiring any skill into the Linux Mint Phase KubesTYLES agent build (ports 8009-8016).

Key Findings

1. **Canonical SKILL.md content recovered for 3/4.** `algorithmic-art`, `mcp-builder`, and `skill-creator` are fully verifiable on

`github.com/anthropics/skills/tree/main/skills/<name>/SKILL.md` and mirror sites (skills.sh, mcpservers.org, aibuilder.sh). `mobius-vocabulary` returns zero hits across GitHub, agentskills.io, and Anthropic mirrors — confirming it is local to `/mnt/skills/user/`, not `/mnt/skills/examples/`.

2. **Three-tier progressive disclosure is universal.** Every skill obeys the (metadata ≤100 words) → (SKILL.md body ≤500 lines) → (bundled resources, on-demand) loading order. github This aligns natively with Möbii³us INNER/INTRA/INTER tier law: metadata = INNER (always loaded), body = INTRA (loaded on trigger), references = INTER (loaded on need).
3. **One hard brand-clamp conflict.** `algorithmic-art` enforces Anthropic Poppins/Lora typography skills MCP Servers and the `#141413/#faf9f5/#b0aea5` palette as **FIXED**. Aibuilder This collides with any Möbii³us TOFAR-stamped output that requires DoinBetter Collective brand tokens — **Guardian gate review required** before any RENDER/VISUALIZE output ships externally.
4. **Evaluation gates differ in shape.** `skill-creator` requires N=3 runs/configuration with mean±stddev benchmark.json; `mcp-builder` requires exactly 10 read-only/independent/idempotent QA pairs in XML. GitHub The two are **complementary**, not competing — `mcp-builder` evaluates the server, `skill-creator` evaluates the skill that uses it.
5. **The "pushy description" doctrine.** `skill-creator` explicitly instructs that descriptions be over-specified to combat *under*-triggering. GitHub Smithery This is in tension with Möbii³us terse declarative voice; the audit recommends keeping descriptions pushy at the YAML layer while keeping body prose terse.

Details

Below is the complete, copy-paste-ready, mkdocs-compatible repository structure. Every block is a separate file; tree is preserved.

📂 Repository Tree

```
doinbetter/
└── audits/
    └── skills/
        └── 2026-05-skill-quad-parse/
            ├── README.md                          ← top-
level repo manifest
            ├── _index/
            │   ├── cross-reference.md             ←
shared concepts/vocabulary
            │   ├── grammatical-diagrams.md        ←
sentence-tree per skill
            │   └── lattice-integration.md         ← 42-
domain + 8-verb wiring
            ├── algorithmic-art/
            │   ├── shares.md
            │   ├── specs.md
            │   ├── questions.md
            │   └── commentary.md
            ├── mcp-builder/
            │   ├── shares.md
            │   ├── specs.md
            │   ├── questions.md
            │   └── commentary.md
            ├── mobius-vocabulary/
            │   ├── shares.md
            │   ├── specs.md
            │   ├── questions.md
            │   └── commentary.md
            └── skill-creator/
                ├── shares.md
                ├── specs.md
                ├── questions.md
                └── commentary.md
```

📄 `README.md` — Top-level Repo Manifest

```
# Skill Quad Parse — 2026-05
```

```
**Audit author:** David Daniel ("Buddy") · DoinBetter Collective · Saint Paul, MN
**Scope:** Four skills: `algorithmic-art`, `mcp-builder`, `mobius-vocabulary`, `skill-creator`
**Purpose:** Pre-flight integration audit before Linux Mint Phase KubesTYLES agent build
**Tier:** INTRA (audit artifact) → promotes to INNER on Guardian sign-off
**TOFAR envelope:** pending hash on commit

## Manifest

| Skill | Source | Tier | Verb-domain | Port |
|---|---|---|---|---|
| algorithmic-art | anthropics/skills (Apache-2.0) | INTRA | RENDER, VISUALIZE | 8014 |
| mcp-builder | anthropics/skills (Apache-2.0) | INTRA | SPEAK, LISTEN | 8011, 8012 |
| mobius-vocabulary | DoinBetter Collective (private) | INNER | GOVERN, AUTHENTICATE | 8009, 8013 |
| skill-creator | anthropics/skills (Apache-2.0) | INTRA | MINT, COORDINATE | 8010, 8015 |

## Read order

1. `_index/grammatical-diagrams.md` — get the shape of each skill in one page
2. `_index/cross-reference.md` — see what they share
3. Per-skill `shares.md` → `specs.md` → `questions.md` → `commentary.md`
4. `_index/lattice-integration.md` — wire to the 42-domain mesh

## Invariants honored

- α = 0.743 (trust floor) — no skill ships below this
- 60/30/10 CRED splits — applied to skill-creator eval workspaces
- ψ = 42 (mesh density) — preserved across the 8-verb fan-out
- Three-tier law (INNER/INTRA/INTER) — immutable
- Branch-free LUT logic — preferred where skills permit
```

```
- ACK-not-BLOCK ethics routing — Guardian gates ACK
first, escalate second
```

📄 `_index/grammatical-diagrams.md` — Sentence-Tree per Skill

```
# Grammatical Diagrams

Each skill rendered as a Reed-Kellogg-style diagram.
Subject is the skill;
Predicate is the dominant verb; Direct Object is what it
produces;
Indirect Object is the recipient; Modifiers constrain;
Subordinate Clauses
nest dependencies.

---

## 1. algorithmic-art

    algorithmic-art ——— creates ——— (philosophy.md +
viewer.html)
         │                 │              │
         │                 │              └─→ to:
claude.ai artifact / browser
         │                 │
         │                 └── modifiers:
         │                       • via p5.js@1.7.0 (CDN)
         │                       • with seeded
randomness
         │                       • under Anthropic brand
clamp (Poppins/Lora)
         │                       • without copying
living artists
         │
         └── subordinate clauses:
               ├── that reads templates/viewer.html
FIRST (STEP 0)
               ├── that fixes seed → randomSeed() +
noiseSeed()
```

```
               └── that emphasizes craftsmanship
REPEATEDLY

---

## 2. mcp-builder

    mcp-builder ─── guides ─── (MCP server scaffold)
         │             │             │
         │             │             └─→ to: external-
API integration agent
         │             │
         │             └── modifiers:
         │                   • in TypeScript (preferred)
or Python (FastMCP)
         │                   • over Streamable HTTP
(remote) or stdio (local)
         │                   • with stateless JSON
         │                   • through 4 phases
(Research/Build/Test/Eval)
         │
         └── subordinate clauses:
               ├── that names tools with service
prefixes (e.g. github_create_issue)
               ├── that annotates
readOnlyHint/destructiveHint/idempotentHint/openWorldHin
t
               ├── that paginates and returns actionable
errors
               └── that produces 10 independent read-
only QA pairs in XML

---

## 3. mobius-vocabulary

    mobius-vocabulary ─── translates ─── (term ↔
operator ↔ geometry)
         │                    │               │
         │                    │               └─→
to: every other skill,
         │                    │
Buddy/TinyBox runtime,
```

```
         │                    │
Guardian gate
         │                    │
         │                    └── modifiers:
         │                          • across quantum /
mathematical / mythic / technical
         │                          • via Π Δ ⊕ᵥ Ω Λ Σ
operator set
         │                          • under
"vocabularies under guard" pattern
         │                          • inside α=0.743
trust floor
         │
         └── subordinate clauses:
               ├── that resolves TOFAR / CRED / Init.0 /
sedenion lexemes
               ├── that maps INNER → INTRA → INTER
tier law
               └── that gates terms before they cross
domain boundaries

---

## 4. skill-creator

    skill-creator ─── mints ─── (new SKILL.md + evals +
benchmark.json)
         │              │              │
         │              │              └─→ to: skill
author + Claude runtime
         │              │
         │              └── modifiers:
         │                    • through draft → test →
review → improve loop
         │                    • with N=3
runs/configuration variance analysis
         │                    • via with-skill /
without-skill paired baselines
         │                    • optimizing description
for triggering accuracy
         │
         └── subordinate clauses:
               ├── that scaffolds skill-
```

```
name/{SKILL.md,scripts/,references/,assets/}
               ├── that generates 20 trigger-eval
queries (8-10 +, 8-10 −)
               ├── that runs 60/40 train/held-out split,
max 5 iterations
               └── that packages via
scripts/package_skill.py to .skill bundle
```

📄 `_index/cross-reference.md` — Shared Concepts

```
# Cross-Reference Index

## Concept-share matrix

| Concept / Vocabulary | algorithmic-art | mcp-builder |
mobius-vocabulary | skill-creator |
|---|:-:|:-:|:-:|:-:|
| YAML frontmatter (name, description) | ✓ | ✓ | ✓ | ✓
(defines it) |
| Progressive disclosure (3-tier) | ✓ | ✓ | ✓ | ✓
(defines it) |
| Reference files in `references/` | — | ✓
(mcp_best_practices, evaluation,
python_/node_mcp_server) | ✓ (vocabulary tables) | ✓
(schemas.md) |
| Templates in `templates/` or `assets/` | ✓
(viewer.html, generator_template.js) | — | ✓ (operator
glyphs) | ✓ (eval_review.html) |
| Scripts directory | — | — | — | ✓ (package_skill,
run_loop, generate_review, aggregate_benchmark) |
| Seeded determinism | ✓ (randomSeed/noiseSeed) | — | ✓
(TOFAR hash seed) | ✓ (eval reproducibility) |
| Pushy description doctrine | inherits | inherits |
should adopt | **defines** |
| XML evaluation format | — | ✓ (qa_pair) | — | inherits
indirectly |
| "Vocabularies under guard" | brand clamp | tool-prefix
discipline | **defines** | description vocabulary |
| TOFAR-hash envelope | output stamp | API auth payload
| **mints** | benchmark stamp |
```

```
| CRED 60/30/10 splits | — | — | governs | applied to
eval train/test |
| Guardian gate review | brand override | tool
annotation review | **enforces** | description sign-off
|

## Vocabulary that must round-trip cleanly

The mobius-vocabulary skill MUST be able to translate
the following Anthropic-native terms into Möbii³us-
native terms without information loss:

| Anthropic term | Möbii³us mapping (proposed) |
|---|---|
| SKILL.md | INTRA-tier manifest leaf |
| frontmatter description | trigger-vocabulary under
guard |
| progressive disclosure | INNER/INTRA/INTER tier law |
| bundled resource | sedenion sub-leaf |
| eval / assertion | Guardian gate predicate |
| tool annotation (readOnlyHint…) | ACK-not-BLOCK
predicate |
| benchmark.json | TOFAR-hashed envelope |
| seed | Init.0 |
| with_skill / without_skill | Π / Δ paired branch |
| feedback.json | ⊕ᵥ (verb-domain merge) |

## Conflict zones (require Guardian gate)

1. **Brand clamp** — algorithmic-art FIXES Anthropic
Poppins/Lora + #141413/#faf9f5; DoinBetter Collective
brand may differ. Resolve before RENDER ships
externally.
2. **Voice tension** — skill-creator's "pushy
description" vs. Möbii³us terse declarative voice.
Resolution: pushy in YAML, terse in body.
3. **License mixing** — three skills are Apache-2.0;
mobius-vocabulary is private. Document downstream
redistribution rules.
```

📄 `_index/lattice-integration.md`

```
# 42-Domain Lattice + 8 Verb-Domain Integration

## Verb-domain port assignments (proposed)

| Port | Verb | Skill responsible | Notes |
|:-:|---|---|---|
| 8009 | GOVERN | mobius-vocabulary | tier-law
enforcement, Guardian gate |
| 8010 | COORDINATE | skill-creator | iteration loop,
test orchestration |
| 8011 | LISTEN | mcp-builder (server) | inbound MCP
requests, stdio + HTTP |
| 8012 | SPEAK | mcp-builder (client) | outbound API
calls, paginated |
| 8013 | AUTHENTICATE | mobius-vocabulary | term
legitimacy, TOFAR verify |
| 8014 | VISUALIZE | algorithmic-art | p5.js artifact
rendering |
| 8015 | MINT | skill-creator | new .skill bundle
creation |
| 8016 | RENDER | algorithmic-art | HTML artifact
output, brand-clamped |

## 42-domain mesh (ψ=42) wiring

The 42 domains fan out from the 8 verb-domains as 8×5 +
2 reserve. Each
skill claims its primary verb at α=0.743 trust floor;
secondary verbs require
⊕ᵥ merge through mobius-vocabulary as broker.

    mobius-vocabulary (broker, INNER)
        ├── governs ──→ all 42 domains
        ├── authenticates ──→ all cross-domain TOFAR
envelopes
        └── routes (ACK-not-BLOCK) ──→ Guardian gate on
each transition

    skill-creator (forge, INTRA)
        ├── mints ──→ new skills into INTRA
        └── coordinates ──→ test/baseline subagents (or
serial in claude.ai)
```

```
    mcp-builder (transport, INTRA)
        ├── listens (8011) ──→ external clients
        ├── speaks (8012) ──→ external services
        └── enforces stateless JSON over Streamable HTTP
for remote

    algorithmic-art (artisan, INTRA)
        ├── renders (8016) ──→ self-contained HTML
        └── visualizes (8014) ──→ p5.js seed-navigable
canvas

## Init.0 boot sequence

1. mobius-vocabulary loads first (INNER) → α-floor
sanity check
2. skill-creator loads second (INTRA) → ready to mint
3. mcp-builder loads third (INTRA) → transport up
4. algorithmic-art loads last (INTRA) → renderer warm

All four are LUT-resolvable; no branching required at
boot.
```

📁 `algorithmic-art/`

`shares.md`

```
# algorithmic-art :: SHARES

## Subject
A guide-skill that creates p5.js generative art with
seeded randomness.

## What it exports / contributes to the system

### Verbs claimed
- **RENDER** (port 8016) — outputs single-file self-
contained HTML
- **VISUALIZE** (port 8014) — provides interactive
seed/parameter UI
```

```
### Reusable components
- `templates/viewer.html` — fixed-shell HTML scaffold
(header, sidebar, canvas)
- `templates/generator_template.js` — p5.js best-
practice reference (NOT a pattern menu) [MCP Servers]
(https://mcpservers.org/agent-
skills/anthropic/algorithmic-art)
- Seed-navigation UI primitive: Prev / Next / Random /
Jump-to-seed
- Action-bar primitive: Regenerate / Reset / Download
PNG
- "Algorithmic philosophy" pattern — a 4-6 paragraph .md
manifesto preceding code

### Vocabulary terms exported (consumable by mobius-
vocabulary)
- `algorithmic philosophy` — computational aesthetic
movement [PIA](https://www.piax.org/skills/anthropics-
skills/algorithmic-art) [MCP Servers]
(https://mcpservers.org/agent-
skills/anthropic/algorithmic-art)
- `controlled chaos` — random variation within strict
boundaries [SkillsMP]
(https://skillsmp.com/skills/davekilleen-dex-claude-
skills-anthropic-algorithmic-art-skill-md)
- `seeded variation` — same seed → identical output
(Init.0 semantic)
- `conceptual seed` — subtle niche reference embedded
invisibly in parameters [skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
- `craftsmanship` — repeated exhortation that the
algorithm "took countless hours" [SkillsMP]
(https://skillsmp.com/skills/anthropics-skills-skills-
algorithmic-art-skill-md)

### File paths it produces
- `<work>/<movement-name>.md` — philosophy
- `<work>/<movement-name>.html` — single-file artifact
(p5.js inline from CDN)

### Patterns reusable elsewhere
- **Read-the-template-FIRST** discipline (STEP 0) — also
```

```
applies to skill-creator
  and any future Möbii³us scaffold
- **Embed-everything-inline** for portability
- **CDN-only external dep** (cdnjs `p5.js@1.7.0`) —
keeps ψ-mesh density preserved

## Indirect object (recipient)
claude.ai artifact pane · any modern browser · Buddy
gallery archive
```

`specs.md`

```
# algorithmic-art :: SPECS

## Dependencies
- p5.js v1.7.0 from
`https://cdnjs.cloudflare.com/ajax/libs/p5.js/1.7.0/p5.m
in.js` [skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
- Google Fonts: Poppins (400/500/600), Lora (400/500)

## File structure
    algorithmic-art/
    ├── SKILL.md
    └── templates/
        ├── viewer.html              ← REQUIRED
STARTING POINT
        └── generator_template.js    ← reference only

## Frontmatter contract
    name: algorithmic-art
    description: Creating algorithmic art using p5.js
with seeded randomness
      and interactive parameter exploration. Use this
when users request creating
      art using code, generative art, algorithmic art,
flow fields, or particle
      systems. Create original algorithmic art rather
than copying existing
      artists' work to avoid copyright violations. [MCP
Servers +2](https://mcpservers.org/agent-
skills/anthropic/algorithmic-art)
```

```
    license: Complete terms in LICENSE.txt   (Apache-
2.0)

## Output contract
- Single self-contained HTML file
- Default canvas: `createCanvas(1200, 1200)` [skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
- Seed default: `12345`
- Sidebar order INVARIANT: Seed → Parameters → Colors?
→ Actions [Skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
[SkillsMP](https://skillsmp.com/skills/anthropics-
skills-skills-algorithmic-art-skill-md)

## Brand-clamp invariants (FIXED — Möbii³us must respect
or override via Guardian)
    :root {
      --anthropic-dark:       #141413;
      --anthropic-light:      #faf9f5;
      --anthropic-mid-gray:   #b0aea5;
      --anthropic-light-gray: (truncated in source)
    }
    fonts: Poppins (UI), Lora (display)

## Validation rules
- Same seed MUST always produce identical output
[SkillsMP](https://skillsmp.com/skills/anthropics-
skills-skills-algorithmic-art-skill-md) [PIA]
(https://www.piax.org/skills/anthropics-
skills/algorithmic-art)
- All parameters MUST have UI controls [Aibuilder]
(https://www.aibuilder.sh/skills/anthropic/algorithmic-
art) [PIA](https://www.piax.org/skills/anthropics-
skills/algorithmic-art)
- Seed Prev/Next/Random/Jump MUST work [PIA]
(https://www.piax.org/skills/anthropics-
skills/algorithmic-art)
- Regenerate / Reset / Download PNG MUST work [PIA]
(https://www.piax.org/skills/anthropics-
skills/algorithmic-art)
- No external files beyond p5.js CDN
- No copying existing artists' work (copyright safety)
[PIA](https://www.piax.org/skills/anthropics-
```

```
skills/algorithmic-art)

## Möbii³us integration points
- **TOFAR envelope:** stamp the philosophy.md + final
HTML hash before publish
- **Guardian gate:** required if brand-clamp is
overridden for DoinBetter output
- **CRED split:** 60% craftsmanship narrative / 30%
algorithm / 10% UI per philosophy
- **α=0.743:** craftsmanship language must clear floor
or kick to revision
```

`questions.md`

```
# algorithmic-art :: QUESTIONS

1. **Brand override path?** Skill says "Keep Anthropic
colors/fonts" as FIXED. [Skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
   For DoinBetter Collective external delivery, do we
patch viewer.html locally,
   or fork the skill into `mobius-algorithmic-art`?
Guardian gate decision.

2. **Sedenion mapping?** Each generative algorithm has
16-ish parameters max in
   practice; sedenion has 16 components. Should params
be auto-mapped 1-to-1
   onto sedenion basis vectors so that mobius-vocabulary
can speak about them
   operator-style?

3. **Init.0 vs. user seed.** Skill defaults seed=12345.
Should Möbii³us inject
   Init.0 hash as the boot seed so every artwork is
provably descended from
   a single root?

4. **CDN dependency under quarantine.** p5.js loads from
cdnjs at runtime.
   The Linux Mint Phase KubesTYLES agent build will run
in environments where
```

```
   egress may be policy-restricted. Vendor p5.min.js
into `assets/`?

5. **Animation determinism.** `function draw()` can be
`noLoop()` (static) or
   animated. Animated outputs are not seed-deterministic
across machines due to
   frame-rate drift. Spec requires same-seed-same-output
— does that bar
   animated work, or do we accept "deterministic-up-to-
noLoop"?

6. **Conflict with `mcp-builder`.** If a future MCP
server exposes art
   generation, who owns the philosophy.md authoring step
— this skill or
   the MCP tool's prompt? Decision needed before any
art-via-MCP route lands.

7. **Conceptual seed and PII.** Skill says "subtle niche
reference embedded [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/a
lgorithmic-art/SKILL.md)
   within the algorithm." If the reference encodes user-
private context, do we
   need a TOFAR-hash audit on the philosophy.md before
publish?
```

`commentary.md`

```
# algorithmic-art :: COMMENTARY

## Design rationale (encoded in the skill)

- **Philosophy precedes code.** The skill forces a 4-6
paragraph manifesto
  before any HTML is written. [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/a
lgorithmic-art/SKILL.md) This is the "Philosophy Before
Implementation" [ClaudeWorld](https://claude-
world.com/articles/anthropic-official-skills-complete-
guide/)
```

```
  anti-AI-slop pattern shared across Anthropic creative
skills.
- **Template is the literal starting point.** Not
inspiration — the actual
  bytes. STEP 0 is in caps with a triple-warning emoji.
[Aibuilder]
(https://www.aibuilder.sh/skills/anthropic/algorithmic-
art) This protects the
  user-experience invariants (seed nav, action bar) from
drift across
  thousands of generations.
- **Process over product.** Beauty lives in the
algorithm's execution, not
  any one frame. Each seed is a print from the same
plate. [PIA](https://www.piax.org/skills/anthropics-
skills/algorithmic-art)

## Best practices encoded

- Always `randomSeed(seed)` AND `noiseSeed(seed)`
[Skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
(both, not one) [skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
[SkillsMP](https://skillsmp.com/skills/davekilleen-dex-
claude-skills-anthropic-algorithmic-art-skill-md)
- Embed everything inline (single-file portability)
- Single CDN dependency only (p5.js)
- Color harmony, not random RGB [SkillsMP]
(https://skillsmp.com/skills/davekilleen-dex-claude-
skills-anthropic-algorithmic-art-skill-md)
- Composition retained even under randomness

## Anti-patterns the skill explicitly forbids

- Creating HTML from scratch [SkillsMP]
(https://skillsmp.com/skills/anthropics-skills-skills-
algorithmic-art-skill-md)
- Inventing custom styling or color schemes [SkillsMP]
(https://skillsmp.com/skills/anthropics-skills-skills-
algorithmic-art-skill-md)
- Using system fonts or dark themes [skills]
(https://skills.sh/anthropics/skills/algorithmic-art)
```

```
[SkillsMP](https://skillsmp.com/skills/anthropics-
skills-skills-algorithmic-art-skill-md)
- Copying living artists' work [PIA]
(https://www.piax.org/skills/anthropics-
skills/algorithmic-art)
- Treating templates/generator_template.js as a "pattern
menu" — it is not [GitHub]
(https://github.com/anthropics/skills/blob/main/algorith
mic-art/SKILL.md) [MCP Servers]
(https://mcpservers.org/agent-
skills/anthropic/algorithmic-art)

## Known limitations

- Hard brand clamp (Anthropic palette) — see
questions.md
- Animation breaks strict same-seed-same-output
guarantee
- p5.js 1.7.0 is pinned; later versions untested by the
skill

## Möbii³us alignment

- **Three-tier law:** SKILL.md = INTRA, philosophy.md =
INTRA-product, the
  final HTML = INTER (ships outward). Clean fit.
- **Branch-free LUT:** seed → noise → particle position
is pure-functional;
  ideal for LUT-style determinism.
- **ACK-not-BLOCK:** the skill never refuses; it absorbs
vague input as
  "creative freedom." This matches Möbii³us ethics
routing — accept input,
  then craft within constraints.
- **Buddy / TinyBox fit:** TinyBox can host the artifact
server-free since
  the HTML is self-contained.

## Complement vs. contrast with the other three

- **vs. mcp-builder:** algorithmic-art produces a UX
endpoint; mcp-builder
  produces a tool endpoint. Composable: an MCP
```

```
`art_generate` tool could
  delegate to this skill.
- **vs. mobius-vocabulary:** algorithmic-art emits
brand-clamped output;
  mobius-vocabulary must register a translator if
DoinBetter brand differs.
- **vs. skill-creator:** algorithmic-art is an exemplar
of the "subjective
  output, no formal evals needed" branch of the skill-
creator decision tree.
```

📁 `mcp-builder/`

`shares.md`

```
# mcp-builder :: SHARES

## Subject
A guide-skill that produces high-quality MCP servers in
TS or Python.

## What it exports / contributes

### Verbs claimed
- **SPEAK** (port 8012) — outbound tool calls to
external APIs
- **LISTEN** (port 8011) — inbound MCP requests over
Streamable HTTP / stdio

### Reusable components
- 4-phase workflow: Research → Implement → Review →
Evaluate [Claudecn](https://claudecn.com/en/blog/claude-
official-skills-walkthrough/) [skills]
(https://skills.sh/anthropics/skills/mcp-builder)
- Tool-naming pattern: `<service>_<verb>_<noun>` (e.g.
`github_create_issue`, [Skills]
(https://skills.sh/anthropics/skills/mcp-builder)
  `slack_send_message`) [GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md) — directly
```

```
applicable to Möbii³us domain naming
- Annotation schema: readOnlyHint / destructiveHint /
idempotentHint /
  openWorldHint [Skills]
(https://skills.sh/anthropics/skills/mcp-builder)
[GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md)
- 10-question XML evaluation harness — pluggable into
any external service
- MCP Inspector workflow: `npx
@modelcontextprotocol/inspector`

### Vocabulary terms exported (consumable by mobius-
vocabulary)
- `tool annotation hint` — non-binding metadata;
semantically near
  ACK-not-BLOCK predicate
- `actionable error message` — error that points the
agent to a fix
- `API coverage vs. workflow tool` — composability
tradeoff
- `stateless JSON` — preferred remote-server transport

### File paths it produces (Python pattern)
    my_mcp_server/
    ├── __init__.py
    ├── server.py
    ├── tools/
    │   ├── __init__.py
    │   └── api_tools.py
    └── requirements.txt [Mintlify]
(https://www.mintlify.com/anthropics/skills/reference/mc
p-builder)

### File paths it produces (TS pattern)
    server with package.json, tsconfig.json,
    Zod input schemas, server.registerTool() calls

## Indirect object
Any external service (GitHub, Slack, Asana, Notion,
Linear, …) — and, by
```

```
extension, every Möbii³us domain that wants to expose
itself as a tool.
```

`specs.md`

```
# mcp-builder :: SPECS

## Dependencies
- **Recommended stack:** TypeScript + Streamable HTTP
(stateless JSON) [Killer-Skills](https://killer-
skills.com/en/skills/anthropics/skills/mcp-builder)
- **Alternate stack:** Python FastMCP (`mcp` package)
- **Schema libs:** Zod (TS) or Pydantic (Python)
- **Inspector:** `@modelcontextprotocol/inspector`

## Reference files (load on demand)
- `reference/mcp_best_practices.md` — naming, response
format, pagination,
  transport, security, error standards [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/m
cp-builder/SKILL.md) [skills]
(https://skills.sh/anthropics/skills/mcp-builder)
- `reference/python_mcp_server.md` — FastMCP patterns
- `reference/node_mcp_server.md` — TypeScript SDK
patterns
- `reference/evaluation.md` — 10-question evaluation
methodology

## Frontmatter contract
    name: mcp-builder
    description: Guide for creating high-quality MCP
(Model Context Protocol)
      servers that enable LLMs to interact with external
services through
      well-designed tools. Use when building MCP servers
to integrate external
      APIs or services, whether in Python (FastMCP) or
Node/TypeScript (MCP SDK). [MCP Servers]
(https://mcpservers.org/agent-skills/anthropic/mcp-
builder) [MCP Servers](https://mcpservers.org/agent-
skills)
    license: Complete terms in LICENSE.txt   (Apache-
```

```
2.0)

## Tool-naming invariants
- MUST include service prefix: `slack_send_message`, NOT
`send_message` [GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md)
- MUST be action-oriented: start with verb (get, list,
search, create, …) [GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md)
- MUST be specific: avoid generic names that could
conflict [GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md) cross-server

## Tool annotation invariants
- Annotations are HINTS, not security guarantees
[GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md)
- Clients MUST NOT make security-critical decisions
based solely on annotations [GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md)

## Evaluation invariants (Phase 4)
- Exactly **10** questions
- Each question MUST be: independent · read-only ·
complex · realistic ·
  verifiable · stable [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/m
cp-builder/reference/evaluation.md) [skills]
(https://skills.sh/anthropics/skills/mcp-builder)
- Format: XML `<evaluation><qa_pair><question>…
</question><answer>…</answer>
  </qa_pair>…</evaluation>`
- Answers MUST be string-comparable [skills]
(https://skills.sh/anthropics/skills/mcp-builder)

## Transport invariants
- Remote: Streamable HTTP, stateless JSON [Mintlify]
(https://www.mintlify.com/anthropics/skills/reference/mc
```

```
p-builder)
- Local: stdio [Killer-Skills](https://killer-
skills.com/en/skills/anthropics/skills/mcp-builder)
[Mintlify]
(https://www.mintlify.com/anthropics/skills/reference/mc
p-builder)
- Authentication (current spec): OAuth 2.1, with RFC
8707 Resource Indicators [DEV Community]
(https://dev.to/x4nent/complete-guide-to-mcp-model-
context-protocol-in-2026-architecture-implementation-
and-4a11)

## External canonical URLs the skill instructs Claude to
fetch
- `https://modelcontextprotocol.io/sitemap.xml` [Killer-
Skills](https://killer-
skills.com/en/skills/anthropics/skills/mcp-builder)
[skills](https://skills.sh/anthropics/skills/mcp-
builder)
-
`https://raw.githubusercontent.com/modelcontextprotocol/
python-sdk/main/README.md` [skills]
(https://skills.sh/anthropics/skills/mcp-builder)
-
`https://raw.githubusercontent.com/modelcontextprotocol/
typescript-sdk/main/README.md` [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/m
cp-builder/SKILL.md) [skills]
(https://skills.sh/anthropics/skills/mcp-builder)

## Möbii³us integration points
- **Each verb-domain port (8009-8016)** is a candidate
MCP server target
- **TOFAR envelope:** wrap each MCP request/response
payload
- **Guardian gate:** annotation review (especially
`destructiveHint=true`)
  triggers human-in-the-loop
- **ψ=42:** if every Möbii³us domain becomes an MCP tool
with consistent
  prefix `mobius_<domain>_<verb>`, the 42-mesh is
preserved
```

`questions.md`

```
# mcp-builder :: QUESTIONS

1. **Stateful vs. stateless under TOFAR.** Skill prefers
stateless JSON [skills]
(https://skills.sh/anthropics/skills/mcp-builder) for
   horizontal scaling. TOFAR envelopes carry session
continuity. Do we layer
   TOFAR over stateless transport, or accept stateful as
a Möbii³us exception?

2. **OAuth 2.1 vs. AUTHENTICATE port (8013).**
Anthropic's MCP spec mandates
   OAuth 2.1 [DEV Community]
(https://dev.to/x4nent/complete-guide-to-mcp-model-
context-protocol-in-2026-architecture-implementation-
and-4a11) for remote servers. Möbii³us AUTHENTICATE is
owned by
   mobius-vocabulary. Do we delegate or proxy?

3. **Annotation policy with ACK-not-BLOCK.**
`destructiveHint=true` does not
   block in MCP — it hints. ACK-not-BLOCK ethics routing
also accepts first.
   Compatible, but who logs the ACK and where does the
audit trail live?

4. **The 10-question floor.** mcp-builder mandates
exactly 10. skill-creator
   suggests 2-3 test prompts then 20 trigger queries. Do
these compose, or do
   we adopt mcp-builder's 10 as Möbii³us minimum for any
tool-exposing skill?

5. **Pagination contract.** Skill recommends pagination
for context
   management. [skills]
(https://skills.sh/anthropics/skills/mcp-builder) Does
mobius-vocabulary need a pagination-vocabulary so that
   `next_page_token` round-trips through Π/Δ predicates
without losing meaning?
```

```
6. **Conflict with `algorithmic-art`.** If we wrap
algorithmic-art as an MCP
   server, the philosophy-first discipline collides with
MCP's "input schema
   + handler" purity. Resolution: a
`mobius_art_generate` tool that takes a
   `philosophy_md` field as required input.

7. **Server prefix.** Should every Möbii³us-emitted MCP
server use prefix
   `mobius_` or `dbc_` (DoinBetter Collective)? Naming
is forever.
```

`commentary.md`

```
# mcp-builder :: COMMENTARY

## Design rationale

- **Quality measured by LLM task completion, not
endpoint count.** This is
  the skill's most important sentence: a server with 5
great tools beats one
  with 50 mediocre ones. [MCP Servers]
(https://mcpservers.org/agent-skills/anthropic/mcp-
builder)
- **API coverage > workflow tools, when uncertain.**
Composability wins
  because LLMs can recompose; [Skills]
(https://skills.sh/anthropics/skills/mcp-builder)
workflows hard-code one path.
- **Clear, consistent prefixes** make agents discover
tools in O(prefix-match),
  not O(read-all-descriptions).
- **Actionable errors** turn a failure into a self-
corrective loop.

## Best practices encoded

- Stateless JSON over Streamable HTTP for remote scale
[Killer-Skills](https://killer-
```

```
skills.com/en/skills/anthropics/skills/mcp-builder)
[Mintlify]
(https://www.mintlify.com/anthropics/skills/reference/mc
p-builder)
- Zod (TS) or Pydantic (Python) for schema rigor
- Output schemas + `structuredContent` (TS) for client
comprehension [skills]
(https://skills.sh/anthropics/skills/mcp-builder)
- Async/await throughout for I/O [skills]
(https://skills.sh/anthropics/skills/mcp-builder)
- 10 read-only complex evaluation questions in XML
[Skills](https://skills.sh/anthropics/skills/mcp-
builder)

## Anti-patterns the skill forbids

- Generic tool names (`send_message`, `create_issue`)
without service prefix
- Vague descriptions that don't precisely match
functionality [GitHub]
(https://github.com/anthropics/skills/blob/main/mcp-
builder/reference/mcp_best_practices.md)
- Treating annotations as security boundaries
- Skipping evaluation

## Known limitations

- The skill defers a lot to external docs
(modelcontextprotocol.io, SDK READMEs)
  — strong network dependency at authoring time
- 10-question evaluation is shallow vs. production
fuzzing
- Stateless JSON sacrifices some agent-flow continuity

## Möbii³us alignment

- **Three-tier law:** SKILL.md = INTRA; reference files
= INTER (load on
  demand); the produced server = a new INTRA artifact.
- **Branch-free LUT:** tool dispatch is pure name-
lookup. Ideal.
- **ACK-not-BLOCK:** annotation hints are advisory, not
blocking — exact
```

```
  semantic match.
- **TOFAR fit:** every tool call is a candidate
envelope.

## Complement vs. contrast

- **vs. algorithmic-art:** mcp-builder is the
SPEAK/LISTEN twin of
  algorithmic-art's RENDER/VISUALIZE.
- **vs. mobius-vocabulary:** mcp-builder needs
vocabulary discipline; the
  `<service>_<verb>_<noun>` rule IS a vocabulary-under-
guard.
- **vs. skill-creator:** the 10-question rule is the MCP
analogue of
  skill-creator's eval set; same family, different
shape.
```

📁 `mobius-vocabulary/`

> *Note: this skill has no public footprint. Specs are reconstructed from the*
> *task description and the canonical Möbii³us lexicon Buddy has used in prior*
> *artifacts. Treat as DRAFT pending Buddy's read-through.*

`shares.md`

```
# mobius-vocabulary :: SHARES

## Subject
A universal-translator skill for Möbii³us system
concepts across quantum,
mathematical, mythic, and technical domains.

## What it exports

### Verbs claimed
- **GOVERN** (port 8009) — tier-law enforcement, term
```

```
legitimacy
- **AUTHENTICATE** (port 8013) — TOFAR-hash verification
of incoming terms

### Reusable components
- **Operator alphabet:** Π, Δ, ⊕ᵥ, Ω, Λ, Σ
  - Π — product / paired-baseline
  - Δ — difference / delta-from-baseline
  - ⊕ᵥ — verb-domain merge (XOR with verb context)
  - Ω — terminal / closure / boundary
  - Λ — lambda / parametric binding
  - Σ — summation / aggregation across mesh
- **Conceptual geometry patterns** — Möbius strip, Klein
bottle, sedenion
  16-basis, ψ=42 mesh
- **Vocabularies-under-guard** pattern — every term has
a Guardian-gate
  predicate before crossing tier boundaries
- **Tier-law mapping table** — INNER / INTRA / INTER
- **Three core constants:** α=0.743 trust floor · CRED
60/30/10 · ψ=42

### Vocabulary terms it OWNS (canonical authority)
- TOFAR (Token-Of-First-Acknowledged-Response, hashed
envelope)
- CRED (Collective Recognition / Equity / Distribution
units)
- Init.0 (root seed; boot-time hash)
- sedenion (16-component non-associative number; param
container)
- Guardian gate (human-in-the-loop predicate)
- Buddy / TinyBox (architecture: client / runtime)
- Möbii³us (the system itself; ³us = third-cycle
iteration)
- DoinBetter Collective (organizational principal)

### Patterns it brokers between the other three skills
- Translates `seed` (algorithmic-art) ↔ `Init.0`
(Möbii³us)
- Translates `tool annotation` (mcp-builder) ↔ ACK-not-
BLOCK predicate
- Translates `assertion / eval` (skill-creator) ↔
Guardian-gate predicate
```

```
- Translates `frontmatter description` (all three) ↔
trigger-vocabulary-under-guard

## Indirect object
Every other skill, every domain (8009-8016), every TOFAR
envelope crosser,
Buddy / TinyBox runtime, and any human or agent reading
mixed-domain
Möbii³us documentation.
```

`specs.md`

```
# mobius-vocabulary :: SPECS

## Tier
INNER — loaded first, always in context, immutable at
runtime

## Frontmatter contract (proposed)
    name: mobius-vocabulary
    description: Universal translator for Möbii³us
system concepts. Enables
      understanding of multi-domain vocabulary
(quantum/mathematical/mythic/
      technical), operator syntax (Π, Δ, ⊕ᵥ, Ω, Λ, Σ),
and conceptual geometry
      patterns. Use when working with Buddy/TinyBox
architecture, parsing
      system documentation with mixed-domain
terminology, translating between
      technical and metaphorical descriptions, or
implementing "vocabularies
      under guard" patterns.
    license: DoinBetter Collective — INNER tier,
internal only

## Constants enforced (immutable)
- α = 0.743 (trust floor; below this → kick to revision)
- CRED splits = 60 / 30 / 10
- ψ = 42 (mesh density; the 42-domain lattice)

## Operator signatures
```

```
- Π(a, b) → paired baseline (with-skill / without-skill
analogue)
- Δ(a, b) → a − b in semantic space
- ⊕ᵥ(verb, payload) → verb-tagged merge into mesh
- Ω(scope) → terminal closure, write barrier
- Λ(name, body) → parametric binding (≈ closure /
lambda)
- Σ(domains) → fan-in aggregation across mesh

## File structure (proposed)
    mobius-vocabulary/
    ├── SKILL.md
    ├── references/
    │   ├── operator-alphabet.md
    │   ├── tier-law.md
    │   ├── tofar-envelope.md
    │   ├── cred-distribution.md
    │   ├── sedenion-basis.md
    │   └── guardian-gate.md
    └── assets/
        └── glyph-table.md

## Validation rules
- Any term crossing INNER → INTRA must be TOFAR-hashed
- Any term crossing INTRA → INTER must clear Guardian
gate
- Any new term added must declare its tier and verb-
domain affinity
- α < 0.743 on any envelope → automatic revision request

## Integration points (the other three skills)
- **algorithmic-art:** brokers `seed`/`Init.0`
translation; provides brand
  override authority
- **mcp-builder:** owns the `mobius_<domain>_<verb>`
prefix scheme; brokers
  annotation ↔ ACK-not-BLOCK
- **skill-creator:** provides Guardian-gate predicate
library for assertions;
  provides trigger-vocabulary review on each new
description
```

`questions.md`

```
# mobius-vocabulary :: QUESTIONS

1. **Where does this skill live on disk?** Task says
   `/mnt/skills/user/mobius-vocabulary/SKILL.md`. The
other three are in
   `/mnt/skills/examples/`. Confirm `user/` path is the
canonical INNER
   location, and that it precedes `examples/` in load
order.

2. **License clarity.** The Anthropic three are Apache-
2.0. mobius-vocabulary
   is INNER / private. If a derived artifact mixes both
(very likely), what's
   the redistribution rule?

3. **Operator collision risk.** Π / Δ / Σ are also
standard mathematical
   notation. Do we need glyph disambiguation in mixed-
context prose?

4. **TOFAR hash algorithm.** SHA-256? BLAKE3? An α-keyed
HMAC? Spec it before
   any real envelope ships, or different domains will
pick different hashes.

5. **CRED 60/30/10 split target.** Across what units?
(Tokens? Time? Compute?
   Attribution credits?) Each interpretation produces a
different audit trail.

6. **ψ=42 vs. 8 verb-domains.** 42 ÷ 8 = 5.25. The mesh
isn't an integer
   fan-out. Are 2 domains held in reserve (as I propose
in
   `_index/lattice-integration.md`), or is the lattice
non-uniform by design?

7. **Guardian-gate latency.** If every cross-tier term
requires human review,
```

```
   we have a throughput problem. Does the gate accept
ACK-now / verify-later,
   or strict-block until human?

8. **Conflict with skill-creator's "pushy
description".** skill-creator
   advises descriptions be pushy to combat under-
triggering. Möbii³us voice
   is terse. Resolution proposed: pushy YAML + terse
body — confirm.

9. **Sedenion params and algorithmic-art.** If we map
p5.js params 1-to-1
   onto sedenion basis, that's an INTRA-tier coupling.
Worth doing, or does
   it over-constrain the artist?

10. **Buddy / TinyBox boundary.** Where does Buddy end
and TinyBox begin?
    Vocabulary should mark the seam clearly so MCP
servers don't straddle.
```

`commentary.md`

```
# mobius-vocabulary :: COMMENTARY

## Design rationale (inferred)

- **Vocabulary is infrastructure.** When four domains
(quantum, math, mythic,
  technical) meet, drift in a single term causes silent
misrouting. A
  universal translator at INNER tier is the cheapest
place to fix this.
- **Operators over names.** Π / Δ / ⊕ᵥ / Ω / Λ / Σ
compress description
  length and remove ambiguity. The cost is glyph
literacy.
- **Vocabularies under guard.** No term enters the mesh
without a Guardian
  predicate. This is the Möbii³us answer to MCP's
`destructiveHint`: not a
```

```
  hint, a gate.

## Best practices encoded (proposed)

- Declare tier and verb-domain for every new term
- TOFAR-hash on every cross-tier transition
- Guardian-gate-then-ACK on every cross-domain
transition
- Reserve 2 of 42 mesh nodes for emergency expansion

## Anti-patterns to avoid

- Inventing operators ad-hoc
- Shipping a term without tier declaration
- Treating CRED splits as approximate
- Letting α drop below 0.743 silently

## Known limitations

- No public peer review (private skill)
- Operator literacy is a learning curve for
collaborators
- Tied to Buddy's voice — succession plan needed

## Möbii³us alignment

- This skill IS the alignment layer. It defines the law
the other three obey.

## Complement vs. contrast

- **vs. algorithmic-art:** mobius-vocabulary GOVERNs;
algorithmic-art RENDERs.
  The vocabulary skill can override art's brand clamp
via Guardian gate.
- **vs. mcp-builder:** mobius-vocabulary AUTHENTICATEs;
mcp-builder SPEAKs/
  LISTENs. Vocabulary owns the prefix scheme; builder
consumes it.
- **vs. skill-creator:** mobius-vocabulary supplies the
Guardian-gate
```

```
  predicate library that skill-creator's assertions can
call into.
```

📁 `skill-creator/`

`shares.md`

```
# skill-creator :: SHARES

## Subject
A meta-skill that mints, edits, evaluates, and
benchmarks other skills.

## What it exports

### Verbs claimed
- **MINT** (port 8015) — produces new SKILL.md +
scaffolding + .skill bundle
- **COORDINATE** (port 8010) — orchestrates with-skill /
without-skill
  subagent runs and grading

### Reusable components / scripts
- `scripts/aggregate_benchmark.py` — produces
`benchmark.json` + `benchmark.md`
  with mean ± stddev per configuration
- `scripts/run_loop.py` — description-optimization loop
(60% train / 40% test,
  max 5 iterations) [Smithery]
(https://smithery.ai/skills/anthropics/skill-creator)
- `scripts/run_eval.py` — single eval runner
- `scripts/package_skill.py` — produces a `.skill`
bundle
- `eval-viewer/generate_review.py` — HTML viewer (server
or `--static` mode)
- `agents/grader.md` — assertion grading subagent prompt
- `agents/comparator.md` — blind A/B comparison subagent
prompt
- `agents/analyzer.md` — analyst pass on benchmark
patterns
```

```
- `assets/eval_review.html` — trigger-eval HTML review
template
- `references/schemas.md` — full JSON schemas for
evals.json,
  eval_metadata.json, grading.json, benchmark.json,
feedback.json

### Vocabulary terms exported
- `progressive disclosure` — the 3-tier loading model
itself
- `pushy description` — over-specified frontmatter
description to combat
  under-triggering
- `with-skill / without-skill` — paired-baseline pattern
(Π / Δ analogue)
- `iteration-N/eval-
ID/{with_skill,without_skill}/outputs` — workspace
layout
- `feedback.json` — user-review side-channel

### File paths it produces
    <skill-name>/SKILL.md
    <skill-name>/{scripts,references,assets}/
    <skill-name>/evals/evals.json
    <skill-name>-workspace/iteration-<N>/eval-
<ID>/{with_skill,without_skill}/
    <skill-name>-workspace/iteration-<N>/benchmark.
{json,md}
    <skill-name>.skill   ← packaged bundle

## Indirect object
Skill authors (Buddy, contributors) · the Claude runtime
that will trigger
the skill · the human reviewer who reads the eval viewer
```

`specs.md`

```
# skill-creator :: SPECS

## Frontmatter contract
    name: skill-creator
    description: Create new skills, modify and improve
```

```
existing skills, and
      measure skill performance. Use when users want to
create a skill from
      scratch, edit, or optimize an existing skill, run
evals to test a skill,
      benchmark skill performance with variance
analysis, or optimize a
      skill's description for better triggering
accuracy. [github]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) [PIA]
(https://www.piax.org/skills/anthropics-skills/skill-
creator)
    license: Complete terms in LICENSE.txt   (Apache-
2.0)

## Skill anatomy (which skill-creator MINTs)
    skill-name/
    ├── SKILL.md (required: YAML frontmatter + Markdown
body)
    ├── scripts/    — executable code, deterministic
tasks
    ├── references/ — docs loaded on demand
    └── assets/     — templates, icons, fonts [Smithery]
(https://smithery.ai/skills/anthropics/skill-creator)

## Progressive disclosure (3 levels, enforced)
1. Metadata (name + description) — always in context
(~100 words)
2. SKILL.md body — in context when triggered (target
<500 lines)
3. Bundled resources — on demand [github]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md)

## Eval JSON schema (evals/evals.json)
    {
      "skill_name": "<name>",
      "evals": [
        {
          "id": 1,
          "prompt": "<user task>",
          "expected_output": "<description>",
```

```
          "files": ["evals/files/sample1.pdf"],
          "expectations": ["…", "…"]
        }
      ]
    } [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/skill-creator/references/schemas.md)

## Grading JSON contract (REQUIRED FIELD NAMES — viewer depends on these)
    {
      "expectations": [
        { "text": "...", "passed": true|false, "evidence": "..." }
      ]
    }
    // NOT: name/met/details — those will break the viewer [github]
(https://github.com/anthropics/skills/blob/main/skills/skill-creator/SKILL.md)

## Benchmark schema (benchmark.json)
- metadata: skill_name, skill_path, executor_model, analyzer_model,
  timestamp, evals_run, runs_per_configuration [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/skill-creator/references/schemas.md)
- runs[]: eval_id, eval_name, configuration, run_number, result {pass_rate,
  passed, failed, total, time_seconds, tokens, tool_calls, errors},
  expectations[], notes[] [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/skill-creator/references/schemas.md)
- run_summary: per-configuration {pass_rate, time_seconds, tokens} as
  {mean, stddev, min, max} [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/skill-creator/references/schemas.md)

## Trigger-eval contract
- 20 queries total: 8-10 should-trigger + 8-10 should-not-trigger
```

```
- Bad: "Format this data"
- Good: realistic, specific, with file paths / company
names / column letters
- Optimization split: 60% train / 40% held-out test, up
to 5 iterations,
  `best_description` selected by **test** score (not
train) to avoid overfit [github]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) [Smithery]
(https://smithery.ai/skills/anthropics/skill-creator)

## Description optimization invariants
- Each query evaluated 3× for reliable trigger rate
[github]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) [Smithery]
(https://smithery.ai/skills/anthropics/skill-creator)
- Uses `claude -p` CLI (Claude Code only — Claude.ai
cannot)

## Workspace layout
    <skill-name>-workspace/
    ├── skill-snapshot/         (only when improving
existing)
    ├── iteration-1/
    │   ├── eval-<descriptive>/
    │   │   ├── with_skill/outputs/     + timing.json +
grading.json
    │   │   └── without_skill/outputs/  (or old_skill/
when improving)
    │   ├── benchmark.json
    │   └── benchmark.md
    ├── iteration-2/ …
    └── feedback.json

## Communication-style invariants
- "evaluation" / "benchmark" — borderline OK
- "JSON" / "assertion" — needs context cues before using
bare [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md)
- Brief inline definitions OK when in doubt
```

```
## Möbii³us integration points
- **MINT** the new skill into INTRA tier; promote to
INNER only via Guardian
- **TOFAR-hash** every benchmark.json for audit
- **CRED 60/30/10** maps cleanly to 60/40 train/test
(10% reserved for
  live A/B held-out)
```

`questions.md`

```
# skill-creator :: QUESTIONS

1. **Pushy vs. terse.** Skill mandates pushy
descriptions. Möbii³us voice is
   terse declarative. Adopt the resolution: pushy YAML,
terse body. Confirm.

2. **Subagents on Linux Mint Phase KubesTYLES.** Skill
notes [X-CMD](https://www.x-
cmd.com/skill/anthropics/skill-creator/) Cowork has
   subagents but Claude.ai does not. [X-CMD]
(https://www.x-cmd.com/skill/anthropics/skill-creator/)
The Mint Phase agent build — does it
   have subagents? If not, we run test prompts in series
(skill explicitly
   permits this).

3. **`generate_review.py` viewer in headless env.** Use
`--static` flag to
   write standalone HTML. [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) Required for any TinyBox runtime.

4. **`feedback.json` ingestion path.** When the viewer
downloads
   `feedback.json`, [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) we need a directory convention so
the next iteration
   auto-picks it up. Propose: `<workspace>/iteration-
<N>/feedback.json`.
```

```
5. **Description optimization needs `claude -p`.** Skill
says this is Claude
   Code only. Mint Phase agent build — does it have
`claude -p`? If not,
   we skip Description Optimization entirely (skill
explicitly permits this).

6. **`package_skill.py` permission model.** Skill warns
installed skill paths
   may be read-only; stage in `/tmp/<skill-name>/`.
[GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) Mint Phase agent must
   have a writable staging dir.

7. **Grading field-name fragility.** Schema requires
`text`/`passed`/
   `evidence`. Any deviation breaks the viewer. Make
this a Möbii³us
   invariant via mobius-vocabulary.

8. **Conflict with mcp-builder's 10-question rule.**
skill-creator suggests
   2-3 test prompts. mcp-builder mandates 10.
Resolution: 10 is the floor
   for tool-exposing skills; 2-3 is the floor for
subjective-output skills.

9. **`evals/evals.json` schema mentions "expectations";
another part says
   "assertions".** Confirm canonical key. Per
`references/schemas.md` it is
   `expectations` (in evals.json) and the grading also
keys on
   `expectations[]`. Use `expectations` everywhere.

10. **Iteration N+1 baseline policy.** When improving an
existing skill,
    baseline can be the original or the previous
iteration. Möbii³us
    should pick one and stick — propose: previous
iteration (Δ=incremental).
```

`commentary.md`

```
# skill-creator :: COMMENTARY

## Design rationale

- **Iteration over revelation.** The skill assumes you
don't know what the
  skill should be on day one. Draft → test → review →
improve → repeat. [X-CMD](https://www.x-
cmd.com/skill/anthropics/skill-creator/)
- **Quantitative AND qualitative.** Benchmarks alone
overfit to the
  verifiable; human feedback alone is unreliable across
scale. Use both.
- **Train/test split.** 60/40 with `best_description`
selected on test
  prevents overfitting [Smithery]
(https://smithery.ai/skills/anthropics/skill-creator)
the description to the train set — a deliberate
  anti-overfit choice that maps to CRED 60/30/10 if we
read 10 as held-out.
- **Why pushy descriptions:** Claude UNDER-triggers more
than it over-
  triggers. [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) The corrective is to over-
specify. Counter-intuitive but
  empirically grounded in the skill's own loop.

## Best practices encoded

- Spawn with-skill AND without-skill in the SAME turn
(parallel)
- Save `timing.json` immediately on subagent completion
(only chance)
- Generate the eval viewer BEFORE evaluating outputs
yourself (capitalized
  exhortation in skill — "GET RESULTS IN FRONT OF HUMAN
ASAP") [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md)
```

```
- Use `generate_review.py`, not custom HTML [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md)
- Kill the viewer server when done (`kill $VIEWER_PID`)
- Generalize from feedback; don't overfit to the user's
specific examples
- Keep prompts lean; remove what isn't pulling weight
- Explain WHY in instructions, not just WHAT (today's
LLMs reward theory-of-mind)

## Anti-patterns the skill explicitly forbids

- Heavy-handed ALL-CAPS MUSTs (yellow flag — reframe
with reasoning)
- Custom HTML for the eval viewer
- Eyeballing assertions when a script can check
programmatically
- Forcing assertions onto subjective outputs (writing
style, art)
- Simple queries like "read file X" as triggering tests
(won't trigger) [GitHub]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) [X-CMD](https://www.x-
cmd.com/skill/anthropics/skill-creator/)
- Skipping the eval viewer step (especially in Cowork —
explicitly called out) [github]
(https://github.com/anthropics/skills/blob/main/skills/s
kill-creator/SKILL.md) [PIA]
(https://www.piax.org/skills/anthropics-skills/skill-
creator)

## Known limitations

- Description optimization requires `claude -p` (Claude
Code only)
- Subagents required for parallel test execution (Cowork
yes, Claude.ai no)
- Browser required for default eval viewer (`--static`
works around it)
- The "pushy description" doctrine produces verbose YAML
that conflicts with
  terse-voice cultures
```

```
## Möbii³us alignment

- **Three-tier law:** the skill itself codifies the 3-
tier load order. Native fit.
- **Branch-free LUT:** triggering is description-match-
then-load. LUT-shaped.
- **CRED 60/30/10:** maps to 60/40 train/test with 10
held out for live A/B.
- **ACK-not-BLOCK:** "be flexible — if user says 'just
vibe,' do that" — ACK ethic.
- **Buddy / TinyBox:** package_skill.py works on any
filesystem (TinyBox-portable).

## Complement vs. contrast

- **vs. algorithmic-art:** algorithmic-art is the kind
of subjective skill
  that skill-creator says doesn't need formal evals.
- **vs. mcp-builder:** mcp-builder's 10-question XML
eval is a complementary
  evaluation shape; skill-creator's evals.json +
grading.json + benchmark.json
  is a more general framework.
- **vs. mobius-vocabulary:** skill-creator's "pushy
description" rubs against
  Möbii³us terse voice — the friction is the most
important thing this audit
  names. Resolution: pushy YAML, terse body.
```

Recommendations

Staged, concrete next steps with the trigger thresholds that change them.

Stage 1 — Commit the audit (today)

1. Drop the entire tree above into `git4emet/doinbetter/audits/skills/2026-05-skill-quad-parse/` exactly as written.
2. Tag commit `audit/skills/quad-parse/v1`.

3. TOFAR-hash the directory and store the envelope at `audits/_envelopes/2026-05-skill-quad-parse.tofar`.
4. **Threshold to advance:** Buddy reads through and signs off; OR opens specific revisions on the open questions.

Stage 2 — Resolve the three blocker questions (this week)

Adopt these as defaults unless overridden:

1. **Brand override (algorithmic-art).** Default: fork `mobius-algorithmic-art` with DBC palette/typography; keep upstream Apache-2.0 attribution. **Change trigger:** if Anthropic Skills updates their template more often than monthly, switch to a runtime patch instead of fork.
2. **Voice tension (skill-creator).** Default: pushy YAML, terse body. **Change trigger:** if observed under-triggering rate on Möbii³us-shaped queries exceeds 15%, push the body too.
3. **Evaluation shape unification.** Default: `mcp-builder`'s 10-question XML for tool-exposing skills; `skill-creator`'s evals.json schema for everything else. **Change trigger:** if a single skill is BOTH (an MCP-tool-exposing skill that needs benchmarking), require both artifacts.

Stage 3 — Wire the verb-domain ports (next week)

Use the port assignments in `_index/lattice-integration.md`:

- 8009 GOVERN / 8013 AUTHENTICATE → mobius-vocabulary
- 8010 COORDINATE / 8015 MINT → skill-creator
- 8011 LISTEN / 8012 SPEAK → mcp-builder
- 8014 VISUALIZE / 8016 RENDER → algorithmic-art

Threshold to advance: all four skills load in Init.0 boot order without α dropping below 0.743.

Stage 4 — Linux Mint Phase KubesTYLES agent build

Only after Stages 1-3 are green:

1. Vendor `p5.js@1.7.0` into `algorithmic-art/assets/p5.min.js` if egress is restricted.
2. Confirm `claude -p` availability for skill-creator's description optimization; if absent, document skip.
3. Confirm subagent availability; if absent, run test prompts in series per skill-creator's permitted fallback.
4. Confirm writable staging at `/tmp/<skill-name>/` for `package_skill.py`.
5. Smoke-test each skill at α=0.743 floor with one TOFAR envelope round-trip.

Threshold to advance to production: all four pass smoke test; all 10 mcp-builder QA pairs (where applicable) verifiable string-comparable; one full skill-creator iteration completes with non-empty `feedback.json`.

Stage 5 — Merge into the 42-domain mesh

Reserve 2 mesh nodes for emergency expansion as documented. Wire the remaining 40 onto 8 verb-domains × 5 each.

Caveats

1. **mobius-vocabulary canonical content was not publicly recoverable.** All four mobius-vocabulary `*.md` files in this audit are RECONSTRUCTED from the user's task description and from canonical Möbii³us terminology used in prior artifacts. They are flagged DRAFT and require Buddy's read-through before promotion to INNER tier. The user authoritatively owns this skill; the audit author does not.
2. **Anthropic skill content is a 2026-05 snapshot.** The `anthropics/skills` repo is actively maintained (as of 2026-04-07 sync, it had 17 top-level skill directories). Pin to a commit SHA in any production reference.
3. **License mixing.** Apache-2.0 (three skills) plus DoinBetter-private (mobius-vocabulary) is fine internally. For any external

publication, route through Guardian gate to confirm redistribution.

4. **Path discrepancy.** Task names paths under `/mnt/skills/examples/` and `/mnt/skills/user/`. The canonical Anthropic repo uses `skills/<name>/`. The mount-paths are local to Buddy's environment and were not directly inspected; specs assume the local mount mirrors the upstream repo plus a `user/` overlay.
5. **Version drift on p5.js, MCP SDK, and frontmatter conventions** is real. The 10-question rule, the `expectations` / `text` / `passed` / `evidence` field names, the `package_skill.py` script names, the `--static` flag — all are from a single 2026-05 snapshot. Re-pin annually.
6. **Operator alphabet (Π Δ Φ_v Ω Λ Σ).** Inferred from the user's task description; signatures proposed in `mobius-vocabulary/specs.md` are best-guess and need Buddy ratification.
7. **The α=0.743 trust floor, CRED 60/30/10 split, and ψ=42 mesh density** are quoted from the task description and treated as canonical Möbii³us invariants throughout this audit. If any of those values has shifted in prior unaudited artifacts, the cross-reference index will need a recalibration pass.
8. **Forward-looking statements clearly flagged:** the proposed port mapping (8009-8016 → 8 verbs), the 40+2 reserve mesh allocation, and the staged Recommendations are *proposals*, not observations — they are decisions to make, not facts on the ground.